

Part 0. Google Colab Setup

Hopefully you're looking at this notebook in Colab!

1. First, make a copy of this notebook to your local drive, so you can edit it.
2. Go ahead and upload the OnionOrNot.csv file from the [assignment zip \(https://www.cc.gatech.edu/classes/AY2022/cs4650_fall/programming/h2_torch.zip\)](https://www.cc.gatech.edu/classes/AY2022/cs4650_fall/programming/h2_torch.zip) in the files panel on the left.
3. Right click in the files panel, and select 'Create New Folder' - call this folder src
4. Upload all the files in the src/ folder from the [assignment zip \(https://www.cc.gatech.edu/classes/AY2022/cs4650_fall/programming/h2_torch.zip\)](https://www.cc.gatech.edu/classes/AY2022/cs4650_fall/programming/h2_torch.zip) to the src/ folder on colab.

NOTE: REMEMBER TO REGULARLY REDOWNLOAD ALL THE FILES IN SRC FROM COLAB.

IF YOU EDIT THE FILES IN COLAB, AND YOU DO NOT REDOWNLOAD THEM, YOU WILL LOSE YOUR WORK!

If you want GPU's, you can always change your instance type to GPU directly in Colab.

Part 1. Loading and Preprocessing Data [10 points]

The following cell loads the OnionOrNot dataset, and tokenizes each data item

```
In [ ]: # DO NOT MODIFY #
import torch
import random
import numpy as np
RANDOM_SEED = 42
torch.manual_seed(RANDOM_SEED)
random.seed(RANDOM_SEED)
np.random.seed(RANDOM_SEED)
# this is how we select a GPU if it's available on your computer.
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
In [ ]: import pandas as pd
from src.preprocess import clean_text
import nltk
from tqdm import tqdm

nltk.download('punkt')
df = pd.read_csv("OnionOrNot.csv")
df["tokenized"] = df["text"].apply(lambda x: nltk.word_tokenize(clean_text(x.lower())))

[nltk_data] Downloading package punkt to /home/andre/nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

Here's what the dataset looks like. You can index into specific rows with pandas, and try to guess some of these yourself :)

```
In [ ]: df.head()
```

```
Out[ ]:
```

		text	label	tokenized
0	Entire Facebook Staff Laughs As Man Tightens P...		1	[entire, facebook, staff, laughs, as, man, tig...
1	Muslim Woman Denied Soda Can for Fear She Coul...		0	[muslim, woman, denied, soda, can, for, fear, ...
2	Bold Move: Hulu Has Announced That They're Gon...		1	[bold, move, :, hulu, has, announced, that, th...
3	Despondent Jeff Bezos Realizes He'll Have To W...		1	[despondent, jeff, bezos, realizes, he, ', ll,...
4	For men looking for great single women, online...		1	[for, men, looking, for, great, single, women,...

```
In [ ]: df.iloc[42]
```

```
Out[ ]: text      Customers continued to wait at drive-thru even...
label          0
tokenized      [customers, continued, to, wait, at, drive-thr...
Name: 42, dtype: object
```

Now that we've loaded this dataset, we need to split the data into train, validation, and test sets. We also need to create a vocab map for words in our Onion dataset, which will map tokens to numbers. This will be useful later, since torch models can only use tensors of sequences of numbers as inputs. **Go to `src/dataset.py`, and fill out `split_train_val_test`, `generate_vocab_map`**

```
In [ ]: ## TODO: complete these methods in src/dataset.py
from src.dataset import split_train_val_test, generate_vocab_map
df = df.sample(frac=1)

train_df, val_df, test_df = split_train_val_test(df, props=[.8, .1, .1])
train_vocab, reverse_vocab = generate_vocab_map(train_df)
```

```
In [ ]: # this line of code will help test your implementation
(len(train_df) / len(df)), (len(val_df) / len(df)), (len(test_df) / len(df))
```

```
Out[ ]: (0.8, 0.1, 0.1)
```

PyTorch has custom Dataset Classes that have very useful extensions. **Go to `src/dataset.py`, and fill out the `HeadlineDataset` class.** Refer to PyTorch documentation on Dataset Classes for help.

```
In [ ]: from src.dataset import HeadlineDataset
        from torch.utils.data import RandomSampler
        #print(train_df)

        train_dataset = HeadlineDataset(train_vocab, train_df)
        val_dataset = HeadlineDataset(train_vocab, val_df)
        test_dataset = HeadlineDataset(train_vocab, test_df)

        # Now that we're wrapping our dataframes in PyTorch datasets, we can
        # make use of PyTorch Random Samplers.
        train_sampler = RandomSampler(train_dataset)
        val_sampler = RandomSampler(val_dataset)
        test_sampler = RandomSampler(test_dataset)
```

We can now use PyTorch DataLoaders to batch our data for us. **Go to src/dataset.py, and fill out `collate_fn`.** Refer to PyTorch documentation on Dataloaders for help.

```
In [ ]: from torch.utils.data import DataLoader
        from src.dataset import collate_fn
        BATCH_SIZE = 16
        train_iterator = DataLoader(train_dataset, batch_size=BATCH_SIZE,
                                   sampler=train_sampler, collate_fn=collate_fn)
        val_iterator = DataLoader(val_dataset, batch_size=BATCH_SIZE, sampler=
                                   val_sampler, collate_fn=collate_fn)
        test_iterator = DataLoader(test_dataset, batch_size=BATCH_SIZE, sampler=
                                   test_sampler, collate_fn=collate_fn)
```

```
In [ ]: ## Use this to test your collate_fn implementation.

## You can look at the shapes of x and y or put print
## statements in collate_fn while running this snippet

for x, y in test_iterator:
    print(x,y)
    break
test_iterator = DataLoader(test_dataset, batch_size=BATCH_SIZE, s
ampler=test_sampler, collate_fn=collate_fn)
```

```

tensor([[ 174, 2903, 9462,    1,  321,  563,   36, 2201,    1, 18
44, 2854,   12,
        223,    1,   33, 7472,    0,    0,    0,    0,    0],
        [ 11, 1085,  828, 6476,  411,   1,   36,   37,   51,   1
58, 3298,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [  1, 2698, 8440,   33,  321,   33,  274,   52,    1,
41,  617, 4959,
        12, 1016,    0,    0,    0,    0,    0,    0,    0],
        [ 25,  882,  289,  552,  301,   61,  181,  593,  289,  27
74, 2775,   36,
        4604, 5445,   36,  113,    1,    1, 1094,   52,    3],
        [3472, 4372, 1331,    1, 2430,   52,   57, 8872,    1,
0,    0,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [4223, 120,   15,  431,    1,  217, 5502,   24,    1,   6
08,   52,    1,
        925,    0,    0,    0,    0,    0,    0,    0,    0],
        [3019,   61, 3210, 2774,    1, 6016, 6635, 4447,    0,
0,    0,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [3123, 4183, 3475,  217, 5632, 8439,   52, 8594, 5089,
21,  605,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [3283,  411,   61, 2461,  158,  476, 3157,  939, 3310,  21
27, 6876,  411,
        33,  343,  327,   21, 7367,  301,    0,    0,    0],
        [ 68,  147, 2389,   33, 7045,  951, 8886,    0,    0,
0,    0,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [1641, 5892,   61, 5897,  976, 8704,  133,  263,    1,
1,   56,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [  1, 2710,    1,  165,    1,    0,    0,    0,    0,
0,    0,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [ 11,   21, 1021, 6181,  619,  148,  149, 3427,  128,  38
79,    0,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [5753, 4651,  321,    1,    1,   21, 1651,    0,    0,
0,    0,    0,
        0,    0,    0,    0,    0,    0,    0,    0,    0],
        [7449, 6726,  770, 5105, 3296,   36,   81,  783,    1,
1, 6870,  783,
        326, 4127,   30,    0,    0,    0,    0,    0,    0],
        [1280,   36, 6891, 3256,    1, 1752, 1010,   24, 2486,  46
71,   30,   31,
        2121,  976,    1, 5917,    1,    0,    0,    0,    0]])
tensor([1., 0., 1., 1., 0., 0., 1., 0., 0., 0., 0., 1., 1., 1.,
1., 1.])

```

Part 2: Modeling [10 pts]

Let's move to modeling, now that we have dataset iterators that batch our data for us. **Go to `src/model.py`, and follow the instructions in the file to create a basic neural network. Then, create your model using the class, and define hyperparameters.**

```
In [ ]: from src.models import ClassificationModel
model = None
### YOUR CODE GOES HERE (1 line of code) ###
model = ClassificationModel(len(train_vocab), embedding_dim=32, hidden_dim = 32, num_layers = 1, bidirectional = True)

# model.to(device)
# #
### YOUR CODE ENDS HERE ###
```

In the following cell, **instantiate the model with some hyperparameters, and select an appropriate loss function and optimizer.**

Hint: we already use sigmoid in our model. What loss functions are available for binary classification? Feel free to look at PyTorch docs for help!

```
In [ ]: from torch.optim import AdamW

criterion, optimizer = None, None
### YOUR CODE GOES HERE ###
criterion, optimizer = torch.nn.BCEWithLogitsLoss(), torch.optim.AdamW(model.parameters(), lr=0.01)
# scheduler = torch.optim.lr_scheduler.StepLR(optimizer, 1, gamma=0.9)

### YOUR CODE ENDS HERE ###
```

Part 3: Training and Evaluation [10 Points]

The final part of this HW involves training the model, and evaluating it at each epoch. **Fill out the train and test loops below.**

```

In [ ]: # returns the total loss calculated from criterion
def train_loop(model, criterion, iterator):
    model.train()
    total_loss = 0

    for x, y in tqdm(iterator):
        optimizer.zero_grad()
        # x = x.to(device)
        # y = y.to(device)
        ### YOUR CODE STARTS HERE (~6 lines of code) ###
        prediction = model(x)
        prediction = torch.squeeze(prediction,0)
        y = y.unsqueeze(1)
        y = y.round()

        loss = criterion(prediction,y)
        total_loss += loss.item()
        loss.backward()
        optimizer.step()
    # scheduler.step()
    ### YOUR CODE ENDS HERE ###
    return total_loss

# returns:
# - true: a Python boolean array of all the ground truth values
#         taken from the dataset iterator
# - pred: a Python boolean array of all model predictions.
def val_loop(model, criterion, iterator):
    true, pred = [], []
    ### YOUR CODE STARTS HERE (~8 lines of code) ###
    for x, y in tqdm(iterator):
        # x = x.to(device)
        # y = y.to(device)
        # print("x",x)
        # print("y",y)

        preds = model(x)
        preds = torch.flatten(preds)
        for i_batch in range(len(y)):
            true.append(y[i_batch])
            pred.append(torch.round(preds[i_batch]))

    ### YOUR CODE ENDS HERE ###
    return true, pred

```

We also need evaluation metrics that tell us how well our model is doing on the validation set at each epoch. **Complete the functions in src/eval.py.**

```
In [ ]: # To test your eval implementation, let's see how well the untrained model does on our dev dataset.
# It should do pretty poorly.
from src.eval_utils import binary_macro_f1, accuracy
true, pred = val_loop(model, criterion, val_iterator)
print(binary_macro_f1(true, pred))
print(accuracy(true, pred))

100%|██████████| 150/150 [00:02<00:00, 69.12it/s]

0.4216917201986171
0.4279166666666667
```

Part 4: Actually training the model [1 point]

Watch your model train :D You should be able to achieve a validation F-1 score of at least .8 if everything went correctly. **Feel free to adjust the number of epochs to prevent overfitting or underfitting.**


```
In [ ]: TOTAL_EPOCHS = 7
        for epoch in range(TOTAL_EPOCHS):
            train_loss = train_loop(model, criterion, train_iterator)
            true, pred = val_loop(model, criterion, val_iterator)
            print(f"EPOCH: {epoch}")
            print(f"TRAIN LOSS: {train_loss}")
            print(f"VAL F-1: {binary_macro_f1(true, pred)}")
            print(f"VAL ACC: {accuracy(true, pred)}")
```

```
100%|██████████| 1200/1200 [01:35<00:00, 12.61it/s]
100%|██████████| 150/150 [00:02<00:00, 73.79it/s]
```

```
EPOCH: 0
TRAIN LOSS: 775.4171956777573
VAL F-1: 0.7974621640060507
VAL ACC: 0.8170833333333334
```

```
100%|██████████| 1200/1200 [01:34<00:00, 12.68it/s]
100%|██████████| 150/150 [00:01<00:00, 79.89it/s]
```

```
EPOCH: 1
TRAIN LOSS: 723.8004207611084
VAL F-1: 0.8273450802293716
VAL ACC: 0.8391666666666666
```

```
100%|██████████| 1200/1200 [01:34<00:00, 12.65it/s]
100%|██████████| 150/150 [00:01<00:00, 85.69it/s]
```

```
EPOCH: 2
TRAIN LOSS: 706.5095884203911
VAL F-1: 0.8122257420283887
VAL ACC: 0.8304166666666667
```

```
100%|██████████| 1200/1200 [01:34<00:00, 12.72it/s]
100%|██████████| 150/150 [00:01<00:00, 81.95it/s]
```

```
EPOCH: 3
TRAIN LOSS: 698.9961122572422
VAL F-1: 0.8362141298875652
VAL ACC: 0.8475
```

```
100%|██████████| 1200/1200 [01:34<00:00, 12.67it/s]
100%|██████████| 150/150 [00:01<00:00, 82.50it/s]
```

```
EPOCH: 4
TRAIN LOSS: 693.1515847146511
VAL F-1: 0.8448174961697538
VAL ACC: 0.8525
```

```
100%|██████████| 1200/1200 [01:39<00:00, 12.03it/s]
100%|██████████| 150/150 [00:01<00:00, 78.85it/s]
```

```
EPOCH: 5
TRAIN LOSS: 689.8829775452614
VAL F-1: 0.8387900771050092
VAL ACC: 0.8495833333333334
```

```
100%|██████████| 1200/1200 [01:42<00:00, 11.65it/s]
100%|██████████| 150/150 [00:01<00:00, 77.04it/s]
```

```
EPOCH: 6
TRAIN LOSS: 686.4265978038311
VAL F-1: 0.8449919692293513
VAL ACC: 0.8545833333333334
```

We can also look at the models performance on the held-out test set, using the same val_loop we wrote earlier.

```
In [ ]: true, pred = val_loop(model, criterion, test_iterator)
        print(f"TEST F-1: {binary_macro_f1(true, pred)}")
        print(f"TEST ACC: {accuracy(true, pred)}")

100%|██████████| 150/150 [00:01<00:00, 77.23it/s]

TEST F-1: 0.8448861905182727
TEST ACC: 0.8570833333333333
```

Part 5: Analysis [5 points]

Answer the following questions:

1. What happens to the vocab size as you change the cutoff in the cell below? Can you explain this in the context of Zipf's Law (https://en.wikipedia.org/wiki/Zipf%27s_law)?

```
In [ ]: tmp_vocab, _ = generate_vocab_map(train_df, cutoff = 0)
print(len(tmp_vocab))
"""
The cutoff discriminates which words are going to be part of our
vocabulary looking at the number of appearances in our training d
ata.
Setting it to a value of 1 means that all the words that appear m
ore than 0 times enter in our vocabulary.
The progression we experiment is:
cutoff len
-----
1      13298
2      9540
3      7612
4      6340
5      5476
6      4825
7      4296
8      3870
9      3590
We see that the number of words in our vocabulary decreases in wh
at it seems a logarithmic progression.
Zipf's Law states that the frequency of any word is inversely pro
portional to the rank of the word in the freq table.
That means the most frequent words are always at the top of the t
able.
With the cutoff, what we are doing is removing the words with the
less frequency. The most frequent words will have very different
frequencies from one
another but the words at the bottom of the table will share their
frequencies and they will be many. By increasing the cutoff, we r
emove each time less and
less words as there will be less words that share their frequenci
es.
"""
```

25387

```
Out[ ]: "\n\nThe cutoff discriminates which words are going to be part of o
ur vocabulary looking at the number of appearances in our trainin
g data.\nSetting it to a value of 1 means that all the words that
appear more than 0 times enter in our vocabulary.\nThe progressio
n we experiment is:\ncutoff len\n----- ---\n1      13298 \n2
9540\n3      7612\n4      6340\n5      5476\n6      4825\n7
4296\n8      3870\n9      3590\nWe see that the number of words
in our vocabulary decreases in what it seems a logarithmic progre
ssion.\nZipf's Law states that the frequency of any word is inver
sely proportional to the rank of the word in the freq table.\nTha
t means the most frequent words are always at the top of the tabl
e.\nWith the cutoff, what we are doing is removing the words with
the less frequency. The most frequent words will have very differ
ent frequencies from one\nanother but the words at the bottom of
the table will share their frequencies and they will be many. By
increasing the cutoff, we remove each time less and\nless words a
s there will be less words that share their frequencies.\n"
```

2. Can you describe what cases the model is getting wrong in the withheld test-set?

To do this, you'll need to create a new `val_train_loop` (`val_train_loop_incorrect`) so it returns incorrect sequences **and** you'll need to decode these sequences back into words. Thankfully, you've already created a map that can convert encoded sequences back to regular English: you will find the `reverse_vocab` variable useful.

```
# i.e. using a reversed map of {"hi": 2, "hello": 3, "UNK": 0}
# we can turn [1, 2, 0] into this => ["hi", "hello", "UNK"]
```

```
In [ ]: # Implement this however you like! It should look very similar to  
val_loop.  
# Pass the test_iterator through this function to look at errors  
in the test set.  
def val_train_loop_incorrect(model, iterator):  
    for x, y in tqdm(iterator):  
        # x.to(device)  
        # y.to(device)  
  
        errors = []  
  
        preds = model(x)  
        preds = torch.flatten(preds)  
  
        for i_batch in range(len(y)):  
            if y[i_batch] != preds[i_batch]:  
                sentence = []  
                for word in range(len(x[i_batch])):  
                    sentence.append(reverse_vocab[x[i_batch][word  
d].item()])  
                errors.append(sentence)  
        return errors  
  
errors = val_train_loop_incorrect(model, test_iterator)  
for sentence in errors:  
    print(sentence)
```

```
0%|          | 0/150 [00:00<?, ?it/s]

['us', 'airways', 'tweets', 'UNK', 'UNK', 'UNK', 'at', 'angry', '
customer', '', '', '', '', '', '', '', '', '', '', '', '', '',
'', '']
['UNK', 'replaces', 'UNK', 'coin', 'UNK', 'with', 'beer', 'taps',
'', 'UNK', 'winners', 'with', 'brew', '', '', '', '', '', '',
'', '', '', '']
['beloved', 'honorary', 'cat', 'mayor', 'in', 'small', 'alaska',
'town', 'dies', 'at', '20', '', '', '', '', '', '', '', '',
'', '', '']
['eagles', 'fans', 'finally', 'sober', 'enough', 'to', 'return',
'to', 'work', '', '', '', '', '', '', '', '', '', '', '',
'', '']
['news', ':', 'finally', ':', 'the', 'indians', 'are', 'replacing
', 'their', 'racist', 'mascot', 'chief', 'UNK', 'with', 'a', 'whi
te', 'woman', 'wearing', 'a', 'native', 'american', 'halloween',
'costume', '']
['these', 'brave', 'teens', 'went', 'UNK', 'for', '3', 'whole', '
days', 'and', 'miraculously', 'survived', '', '', '', '', '',
'', '', '']
['lou', 'UNK', 'blames', 'stanford', '', 's', 'loss', 'to', 'uta
h', 'on', 'UNK', '', '', '', '', '', '', '', '', '']
['mayor', 'rob', 'ford', 'calls', 'in', 'sick', 'on', 'bob', 'UNK
', 'day', '']
['amsterdam', 'to', 'fly', 'rainbow', 'flag', 'for', 'russian', '
president', 'putin', '', 's', 'visit', 'to', 'the', 'capital',
'', '']
['four', 'UNK', 'took', 'cocaine', 'thinking', 'it', 'was', 'UNK
', '']
['for', 'the', 'first', 'time', 'in', 'saudi', 'arabia', ' ', 'wo
men', 'UNK', 'to', 'issue', 'UNK', ' ', ' ', ' ', ' ', ' ', ' ',
'', '']
['boy', 'suspended', 'over', 'UNK', 'gun', 'seeks', 'to', 'clear
', 'school', 'record', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ',
'', '']
[' ', 'nothing', 'would', 'surprise', 'me', 'at', 'this', 'point
', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ', ' ',
'8', 'separate', 'news', 'items', 'today', ' ', '']
['why', 'north', 'korea', '"s', 'capital', 'is', 'the', 'UNK', 's
cience', 'fiction', 'film', 'set', '"', ' ', ' ', ' ', ' ', ' ',
'', '']
```